

WHG Licensing: Audit, Discussion Points & Recommendations

Internal briefing prepared in response to the TreeAlive licensing inquiry

Prepared by: Dr Stephen Gadd, Technical Director · **Date:** 16 June 2026 · **Audience:** WHG team & (potentially) Pitt commercialization office · **Status:** For discussion — not legal advice

Executive summary

A genealogy startup (TreeAlive) asked which licence governs WHG's published place data, and flagged that our documentation cites **CC BY-NC 4.0** while our public-data page cites **CC BY 4.0**. A full sweep of the codebase confirms the inconsistency is **larger than reported**: WHG states a licence to users in at least **nine** distinct places, and they do not agree. The contradictions split along **audience lines** — contributors are told CC BY 4.0, reusers are told CC BY-NC 4.0 — which is where the genuine legal exposure lies. Crucially, much of the data WHG provides is delivered not as file downloads but **live through the Reconciliation API**, where multiple aggregated sources are blended into single results — a harder case with different and more consequential licensing questions, examined in §2.

The single coherent finding: WHG's aggregated index is a *mix* of licences (CC BY, CC BY-NC, ODbL, public domain, and bespoke terms). It cannot honestly be covered by any single blanket licence in either direction. The legally clean position and the internally consistent position are the same: **state no global licence over the aggregate; let each source carry its own licence, and attach that provenance to the data wherever it leaves the platform.**

This is already being addressed. A sitewide remediation — the *Citations, Licences & Credit* programme — was designed and locked on 6 June 2026 and is substantially built and deployed (§5). This inquiry validates that work and tells us which remaining piece to prioritise.

1. Audit — where WHG states a licence, and what it says

Location (audience)	What it states	Commercial use?
<code>dataset_create</code> / <code>dataset_validate</code> — upload acceptance checkbox (contributors)	"Accept CC BY 4.0 Licence"	Allowed
Public-data page <code>public_list</code> (reusers) ×2	"All such public data... CC BY 4.0 "	Allowed
Search results <code>schema.org</code> JSON-LD (machines / aggregators)	<code>licenses/by/4.0/</code>	Allowed
/license/ policy page <code>license.html</code>	" CC BY-NC 4.0 , unless specified otherwise"	Forbidden
Dataset page <code>ds_metadata</code>	"CC BY- NC 4.0"	Forbidden
Live download stream header <code>download_file.py</code>	"Attribution- NonCommercial 4.0"	Forbidden
	"CC-BY 4.0 NC " (<i>not a valid licence name</i>)	Forbidden

Location (audience)	What it states	Commercial use?
Download README + "ready" email <code>tasks.py</code> / <code>download_ready</code>		
DataCite DOI registration <code>doi.py</code>	CC-BY-NC-4.0 (permanent external record)	Forbidden
Download modal <code>licensing_info</code>	Prose says BY & "freely adapted"; link <i>label</i> says BY-NC; link <i>target</i> is BY — self-contradictory in one paragraph	—

Note: these are the *download and static* surfaces. The Reconciliation API — arguably the larger delivery channel — is treated separately in §2 because its licensing questions are different and harder.

2. Delivery channels: downloads vs. the Reconciliation API

WHG delivers data through two very different channels, and they carry different risks:

- **File downloads** hand over a *single* dataset or collection. Each has a known, attachable licence — the problem here (§1) is purely *consistency*, and it is straightforward to fix.
- **The Reconciliation API** answers queries live by *blending* candidate matches — names, alternate names, geolocation, type — drawn from many aggregated sources at once. This is the harder case.

The reconciliation channel raises four distinct issues:

- **Blended delivery without licence context.** Each reconciliation candidate is returned with only its id, name, type, score and match flag — **no per-result licence or source-rights signal**. A consumer (OpenRefine, or TreeAlive calling programmatically) therefore receives results spanning many differently-licensed sources with no machine-readable way to know which is which — and so cannot comply with per-source terms even if willing.
- **Blending can manufacture a derivative.** Where a record merges attributes from several sources, the output is arguably a compilation of them. If any input carries **ShareAlike** terms (e.g. OSM/OHM under ODbL, or a BY-NC-**SA** source), ShareAlike is *copyleft* — it can force the whole blended record under the same terms, "infecting" otherwise-permissive (even public-domain) data mixed into the same record. For an aggregator, **ShareAlike is a more structural problem than NC**.
- **Facts vs. database right — and a US/UK split.** A single name → coordinate is fact (not copyrightable, safe to reuse). But *systematic, repeated* extraction of a substantial part of the index via the API can trigger **UK/EU sui generis database right** even where no copyright subsists — which is exactly what a genealogy app reconciling at scale does. Note the asymmetry: WHG/Pitt and TreeAlive are US-based (no database right), but several upstream contributors' rights arise under UK/EU law, so one API call can be low-risk for the consumer yet breach an upstream contributor's rights.
- **It is also WHG's strongest commercial asset.** The reconciliation / linking *service* — the matching, ranking and cross-source crosswalk — is the value WHG genuinely creates, distinct from the underlying records. So the API both *complicates* the rights picture **and is precisely the layer WHG can legitimately put behind a commercial licence**, even where it cannot license the source records themselves.

3. The two structural problems behind the audit

3.1 Contributors are told CC BY 4.0; reusers are told CC BY-NC 4.0

These are different grants. CC BY 4.0 is **irrevocable**, **explicitly permits commercial use**, and forbids the licensor from layering additional restrictions ("no additional restrictions"). So for everything that arrived through the normal upload flow, WHG may have **less** right to impose NC than it currently advertises — the opposite of the intuitive worry. The genuine NC exposure comes specifically from the **absorbed upstream sources** (e.g. Humphrey's transcripts and other NC gazetteers we have aggregated), not from the contributor pipeline.

3.2 The DOI registrations bake NC into a permanent external record

Every minted DOI registers `CC-BY-NC-4.0` in DataCite. This has the longest tail of any item in the audit because DOI metadata is externally cached and hard to walk back. It should emit each dataset's *actual* licence before the next batch of DOIs is minted — work already specified in the remediation (§5).

4. Points for discussion

- **WHG is simultaneously licensee and licensor.** We are a *licensee* of upstream data (much of it NC) and a *licensor* to parties such as TreeAlive. We cannot grant downstream what we do not hold upstream — so the aggregate index cannot be blanket-authorised for commercial reuse.
- **What WHG can legitimately commercialise.** A dual-licence model (free NC + paid commercial — cf. Edinburgh's AMI Meeting Corpus, as Palak notes) is viable, but only over what WHG owns or may sublicense: the **software** (already BSD-3), the **reconciliation / linkage service** (§2), and **whole contributed gazetteers whose contributors granted commercial rights**. The raw aggregated place records are mostly *not* ours to sell.
- **"Non-commercial" is poorly defined** (Palak's point). NC famously catches uses we want — commercial academic journals, tuition-charging universities — while not reliably stopping the uses we would want to prevent. NC is a weak protective mechanism on which to build a strategy.
- **TreeAlive's "source-aware attribution" posture aligns with our direction.** Sam's David Rumsey-derived template is exactly the per-source model WHG is moving toward; this is a point of genuine alignment to acknowledge in any reply.

On the University of Portsmouth / Humphrey Southall terms

Humphrey's note asserts: (a) Portsmouth holds *some* copyright and the files are "definitely not public domain"; (b) he is content to release them as **CC BY-NC-SA**; and (c) "*BY*" means *I should be included as a co-author on any publication making substantial use.*" Taking these in turn — fairly, but candidly:

Largely sound: "not public domain" is defensible, though not quite as stated. The raw facts (a place-name, a coordinate) are not copyrightable; what *is* protectable is the **compilation** and — under UK law specifically — the **sui generis database right**, which protects the substantial investment he describes (the OCR-scraping program and the major manual student clean-up). Because that work was done by Portsmouth staff and students, the rights holder is most likely the **University, not Humphrey personally** — so any commercial use needs *institutional* sign-off, not just his.

Over-reaching — the co-authorship claim should be challenged. As a matter of Creative Commons interpretation, the "BY" / Attribution requirement means *credit* — naming the creator, linking the source and licence, and noting any changes. It does **not** confer or require **co-authorship**.

Attribution ≠ authorship: co-authorship is governed by authorship norms (a substantive intellectual contribution to *that specific new work*) and/or by contract, not by a licence's attribution clause — and even academic convention does not grant authorship merely for *using* a dataset, however substantially. He may legitimately *negotiate* co-authorship as a separate, bespoke condition of access, but it cannot be derived from "BY", and as a bespoke term it binds only WHG/Portsmouth — it cannot be passed through to downstream users the way a licence can.

Why this matters operationally: a mandatory co-authorship term is in practice *more* restrictive than NC — it cannot be satisfied automatically, does not scale, and **WHG cannot honour it downstream** (we cannot auto-co-author Humphrey onto a TreeAlive feature). And the **ShareAlike** element he offers is the bigger structural issue for blended delivery (§2). Both point to the same need: the platform must be able to record bespoke terms precisely and **gate specific sources out of channels (downloads, reconciliation) where their terms cannot be met** — capability the remediation in §5 provides.

5. Work already underway to address this

This inquiry did not catch WHG unprepared. A unified **Citations, Licences & Credit** design was reviewed and locked on **6 June 2026**, expressly to fix "inconsistent, partly-incorrect licence handling sitewide," and is now substantially implemented. Its relevant components:

- **One licence vocabulary.** A new **licensing** app with a curated **License** table (SPDX ids — **CC-BY-4.0**, **CC-BY-NC-4.0**, **ODbL-1.0**, **ODC-By-1.0**, **CC0-1.0**, ...) carrying explicit *permits-commercial*, *share-alike* and *attribution-required* flags, plus a *custom* flag for bespoke, non-SPDX terms (exactly the Portsmouth case). One vocabulary, one renderer, reused everywhere a licence is named.
- **Truthful, per-source rights.** Rights statements in downloads, DOIs and the UI now come from each object's *actual* licence, never a hard-coded constant. Aggregates spanning multiple sources emit a **list of per-source rights**, never a single collapsed or "most-restrictive" licence. WHG's own curation/aggregation licence is asserted as a clearly-distinguished **overlay alongside** each source's rights — never instead of it.
- **Structured credit (CRediT).** Contributor roles follow the NISO CRediT taxonomy via a through-model that can attach credit to datasets, collections *and* authority gazetteers, with degree (lead/equal/supporting) and a corresponding-author flag. This is the proper home for bespoke acknowledgements — Aucott's clean-up management, Stoner's scraping program, and any negotiated co-authorship condition — recorded as structured, exportable metadata rather than free text.

- **Attribution follows the data out.** API responses gain a namespace-keyed `attribution` block (source name, citation, licence, rights holder, contributors) plus the WHG overlay, via one shared helper reused across search, place-detail and downloads.
- **Capture at source.** The ingestion pipeline is being upgraded so every authority gazetteer's *actual* upstream licence, citation and rights holder is captured at ingest and pushed into the registry — replacing the stale, incorrect global "CC-BY-NC-4.0" assumption with verified per-source terms (an explicit audit of every namespace is in progress).

Status: the foundation — licence model, per-source rights, structured CRediT credit, and the API attribution block — is built and largely **in production**. The remaining pieces are (i) the per-namespace authority-licence audit and ingest capture (in progress on the development branch and the indexing pipeline), and (ii) attaching per-result provenance to *reconciliation* responses (§6, recommendation 3) — the one place the new attribution block is not yet wired in, which this inquiry tells us to prioritise.

6. Recommendations

6.1 Single coherent position to assert everywhere

WHG's own **software** is BSD-3. WHG **editorial content and value-add** default to CC BY-NC 4.0. **Contributed and absorbed datasets each carry their own licence**, shown on the item page and attached to the data wherever it leaves the platform — reuse must respect the source's terms. **There is no single licence over the aggregate.**

6.2 For the reply to Sam / TreeAlive

- Warm "stay tuned" is appropriate; his current free use is fine today regardless of how this resolves.
- Be honest that the published place data is a *mix*; he should not assume commercial rights over the aggregate, nor over bulk reconciliation output.
- Specific CC BY / CC0 sources are safe for any use; a future commercial path would be a negotiated licence over WHG's reconciliation layer, not the raw records.
- Acknowledge his source-aware attribution approach as aligned with WHG's direction.

6.3 For the Pitt commercialization conversation

Frame the question precisely so the meeting is productive:

"Given a mixed corpus (CC BY, CC BY-NC, ODbL/share-alike, public domain, and bespoke terms such as required co-authorship), what can Pitt/WHG lawfully offer commercially — and how should a dual licence be structured around the WHG-owned value-add (the reconciliation/linkage service) rather than the aggregated source data?"

6.4 Capture the user's use case, and gate on it

The per-source licence flags only become *enforceable* once we know what the user intends to do. The recommendation is therefore to ask users to **declare their use case** at the point of access — and to gate resources on that declaration combined with each source's terms:

- **Where to capture it.** At download (a per-download attestation, extending today's "by clicking download you agree" consent) and at **API-token issuance** (declared once when the token is granted, rather than on every request — the reconciliation service is already token-authenticated).
- **Default to non-commercial; offer an override.** The declaration should **default to non-commercial use** — the safe presumption that fits the great majority of WHG's users and grants the full set under NC terms with no friction. A *commercial* use is then an explicit **opt-in / override**: the user actively declares it, which is what triggers both the gating and the commercial-enquiry funnel below. This keeps the common path frictionless while making the consequential choice deliberate and recorded.
- **How it gates.** The platform combines the declaration with the per-source licence flags. The default non-commercial path receives the full set under the appropriate terms; a *commercial* override is served only the commercially-licensed subset (CC BY, CC0, ODC-By, ...) and is excluded from NC and bespoke-restricted sources (e.g. the Portsmouth material). This declaration is the input that drives the channel gating in the table below.
- **Why it is protective.** The declaration creates an **attestation record** evidencing WHG's good-faith enforcement of upstream NC and bespoke terms, and places compliance responsibility on the declaring user — a meaningful shield for a platform redistributing third-party data.
- **It doubles as a commercial funnel.** A *commercial* declaration can route the user to a licensing enquiry (the Pitt-office path) rather than a flat refusal — turning the gate into lead generation. This is precisely the TreeAlive scenario.

6.5 Will NC-gating gut the commercial offering?

A reasonable objection: if we filter the API to commercially-usable sources, are the results still useful? **For the great majority of queries, yes** — because the commercial-incompatible material is a small, specialist minority sitting on top of a backbone that is almost entirely commercial-compatible. The index composition (record counts as configured; licences as currently understood, pending the \$5 per-namespace audit):

Source (namespace)	~Records	Licence	Commercial use
GeoNames <code>gn</code>	12 m+	CC BY 4.0	✓ attribution
Wikidata <code>wd</code>	8 m+	CC0	✓ public domain
OpenStreetMap <code>osm</code>	6 m+	ODbL	✓ but share-alike
Getty TGN <code>tgn</code>	3 m+	ODC-By 1.0	✓ attribution
GB1900 <code>gb</code>	~2.5 m	CC BY-SA	✓ but share-alike
Pleiades <code>pl</code>	37 k+	CC BY 3.0	✓ attribution
Library of Congress <code>loc</code>	names authority	US Gov / no known ©	✓
OpenHistoricalMap <code>ohm</code>	—	ODbL	✓ but share-alike
Natural Earth / PeriodO	—	public domain	✓
	specialist minority	CC BY-NC / bespoke	✗ excluded

Source (namespace)	~Records	Licence	Commercial use
Contributed NC / bespoke sets (CHGIS, Trismegistos, Ottoman NFS, Portsmouth transcripts, ...)			

The four largest sources alone — roughly **29 million records** — are all commercial-compatible, before counting Pleiades, LOC, Natural Earth and the rest. Three factors make the practical degradation gentler still:

- **Places are heavily duplicated across authorities.** Reconciliation returns *candidates*; if the best match sits in an NC source, the same place usually also exists in GeoNames/Wikidata/TGN, so the commercial user still gets a match — not nothing. The real loss is the long tail of places that exist *only* in a restricted contributed set.
- **WHG's own value-add is unaffected.** The matching, ranking and cross-source linkage — the part WHG owns and would be selling — is served regardless of which source records pass the filter.
- **The loss is concentrated, not uniform.** Commercial users lose precisely the specialist niches (detailed Chinese historical units, Ottoman-era places, parish-level British transcripts) — which are also where commercial demand is thinnest. A genealogy product like TreeAlive is overwhelmingly served by the global backbone.

So the realistic risk is not "useless"; it is two deliberate design decisions:

1. **Whether to include share-alike sources in the commercial subset.** OSM, OHM (ODbL) and GB1900 (CC BY-SA) are commercial-OK but copyleft. Including them (with an obligations notice) preserves coverage; excluding them keeps the subset obligation-free but is where a real dent would be felt — British historical coverage (GB1900) and detailed/modern geometry (OSM/OHM).
2. **Accepting the niche loss** of the restricted contributed sets — acceptable, and arguably correct, since those are the contributors whose terms forbid commercial reuse.

What ShareAlike (copyleft) requires of the end-user

ShareAlike licences — **ODbL** (OSM, OHM) and **CC BY-SA** (GB1900) — permit commercial use but attach a *copyleft* obligation that can propagate into the user's own work. A commercial consumer must:

- **Attribute** the source, as with any attribution licence; and
- **Re-share derivatives under the same terms.** Under CC BY-SA, any *adaptation* of the material must itself be licensed CC BY-SA. Under ODbL, if the user builds and publicly uses a *derivative database* (e.g. ingests the place data into their own enhanced place-database and distributes it), that derivative database must be offered under ODbL with its alterations made available.

The practical consequence: a commercial user generally **cannot keep a derived place-database proprietary** if they distribute it, and SA terms can be **incompatible with other licences** in their stack.

A distinction that matters for app-builders: **ODbL's copyleft attaches to the database, not necessarily to the end-product** (its "Produced Work"). An app that merely *displays* OSM-derived data must attribute and keep any modified database open, but is not forced to open-source its application. **CC BY-SA is broader** — it reaches adaptations of the content itself, not just databases. By contrast **ODC-By (TGN) and CC BY are attribution-only — no share-alike** — making them the cleanest sources for a commercial consumer.

The upshot: SA sources are *usable* commercially but carry obligations the consumer must accept knowingly — which is exactly why attaching **per-result provenance** to reconciliation responses (recommendation 3) matters: the consumer needs to see which candidates are share-alike to decide whether to use and comply, or to filter them out.

Before committing to the gated commercial offering, measure two numbers: (i) `record_count` per namespace from the registry, and (ii) each namespace's frequency in a sample of *real* reconciliation responses — what matters is a source's share of actual hits, not its share of the index. The qualitative conclusion (backbone is commercial-safe; restricted material is a specialist minority) is robust regardless, but these two figures quantify the niche loss precisely.

6.6 Documentation & platform changes (priority-ordered)

#	Location / area	Change	Priority
1	DataCite DOI <code>doi.py</code>	Stop hard-coding <code>CC-BY-NC-4.0</code> ; emit the dataset's actual licence(s). Permanent external record — fix before next minting.	Highest
2	Upload acceptance <code>dataset_create</code> / <code>dataset_validate</code>	Legally load-bearing. Decide deliberately what grant we ask contributors for; counsel signs off on this wording specifically.	Highest
3	Reconciliation API responses	Attach per-result provenance + licence (the namespace attribution block already exists — wire it into reconciliation). Enables consumer compliance and lets WHG gate sources whose terms can't be honoured. <i>The remaining piece of the \$5 programme.</i>	High
4	Public-data page <code>public_list</code> ×2	Replace blanket "CC BY 4.0" with "each dataset/collection carries its own licence, shown on its page." Most-read, currently most wrong (the line Sam quoted).	High

#	Location / area	Change	Priority
5	/license/ policy <code>license.html</code>	Keep CC BY-NC 4.0 as default for WHG-created editorial only; carve out (a) datasets = per-source, (b) software = BSD-3.	Med
6	Download modal <code>licensing_info</code> ; README, email & stream header	Fix the self-contradiction; remove the invalid "CC-BY 4.0 NC" string; emit the item's actual licence (field now exists).	Med
7	Search <code>schema.org</code> JSON-LD	Remove the global <code>by/4.0/</code> or make it per-record; machines consume this as a blanket assertion we cannot back.	Med
8	Use-case declaration & per-source gating	Capture the user's declared use case at download and at API-token issuance — defaulting to non-commercial, with an explicit commercial override . Combine with the per-source flags to serve only the permitted subset and exclude sources whose terms can't be met (§6.4). Records an attestation; doubles as a commercial-enquiry funnel.	High
9	Authority licence audit + default label <code>settings.py</code>	Complete the per-namespace upstream-licence audit; align the default label with the agreed position. <i>In progress</i> .	Underway

Items 4–7 are text/template only; 1, 3 and 6 are small wiring changes that read the per-source licence/attribution data the remediation already provides. The recommended sequence is to settle the *wording* first (it is a legal decision), then a single source-of-truth change makes the surfaces consistent at once.

This document is an internal technical and editorial audit prepared to inform discussion. It identifies factual inconsistencies and structural logic; it is not legal advice. Statements about copyright, database right and Creative Commons interpretation are offered in good faith to frame the questions, not to resolve them. Decisions on contributor grant wording, commercial licensing, third-party (e.g. University of Portsmouth) rights, and any change to DOI metadata should be confirmed with Pitt's Office of Innovation & Entrepreneurship / legal counsel.