

WHG Licensing: Briefing for the Office of Innovation

How the World Historical Gazetteer handles rights over aggregated place data — current state, working evidence, and the questions on which we seek guidance

Prepared for: Tara Bransted, Licensing Manager, Office of Innovation, University of Pittsburgh · **By:** Dr Stephen Gadd, Technical Director, World Historical Gazetteer · **Date:** 29 July 2026 · **Status:** For discussion — not legal advice

1. Summary, and what we are asking for

The World Historical Gazetteer (WHG) is a Pitt-hosted platform that aggregates and links historical place data. Its index currently holds **50.5 million place records** drawn from 26 distinct source gazetteers, plus several hundred contributed datasets. Those sources arrive under **at least nine different licences** — public domain, CC BY, CC BY-SA, ODbL, ODC-By and CC BY-NC — and for a small specialist remainder the terms are not yet recorded at all (§5).

The position we have adopted: WHG asserts *no single blanket licence* over the aggregate. Each source carries its own licence; that licence travels with the data wherever it is surfaced or leaves the platform; and WHG's own curation/aggregation licence (CC BY-NC-4.0) is asserted as an **overlay alongside** each source's terms, never instead of them.

Since an internal audit in June 2026 this has been substantially **built and deployed** — a controlled licence vocabulary, per-source licence capture at ingest, a contributor-facing licence picker, and licence badges surfaced throughout the public interface. §4 shows this working on the live site.

What we would value from this meeting

1. **Confirmation (or correction) of the per-source model** as the right legal posture for an aggregator that is simultaneously licensee and licensor.
2. **Guidance on what Pitt/WHG could lawfully offer commercially**, given that the raw source records are largely not ours to sell (§6).
3. **A view on the contributor grant wording** we ask uploaders to accept — this is the legally load-bearing text and we would like it reviewed (§7.2).
4. **A steer on two specific third-party matters:** a bespoke co-authorship condition attached to one source, and whether a declared-use-case attestation is a meaningful protection (§7.1, §7.4).

2. Why an aggregator is a distinctive case

WHG is not principally a producer of data. It is a **licensee** of many upstream sources and a **licensor** to downstream users, and it cannot grant downstream what it does not hold upstream. Three consequences follow, and they shape everything in §4:

- **No honest blanket licence exists.** A single site-wide licence would either over-claim (purporting to license third-party data we do not own) or under-claim (imposing NC restrictions on CC BY material where "no additional restrictions" applies). Both are wrong; only per-source terms are defensible.

- **ShareAlike matters more than NonCommercial.** Non-commercial terms simply exclude a source from a commercial channel. ShareAlike (ODbL, CC BY-SA) is *copyleft*: where records are blended, it can propagate obligations into the combined product. For an aggregator this is the more structural risk — and, as §5 shows, it affects a much larger share of our corpus than NC does.
- **Facts vs. database right, with a US/UK split.** A single name-to-coordinate pairing is a fact and is not copyrightable. But *systematic* extraction of a substantial part of a compilation can engage the **UK/EU sui generis database right**, which has no US equivalent. WHG/Pitt are US-based; several of our upstream contributors' rights arise under UK law. One bulk API consumer can therefore be low-risk domestically yet still breach an upstream contributor's rights.

3. The model we implemented

Four components, all now in production:

- **One controlled licence vocabulary.** A curated table of 15 licences keyed on SPDX identifiers ([CC-BY-4.0](#) , [ODbL-1.0](#) , [CC0-1.0](#) , ...), each carrying explicit *permits-commercial*, *share-alike*, *attribution-required* and *no-derivatives* flags, plus a *custom* flag for bespoke non-SPDX terms. One vocabulary, one renderer, used everywhere a licence is named — so the surfaces cannot drift apart again.
- **Per-source capture at ingest.** Every authority gazetteer's actual upstream licence, citation, rights holder and source URL is recorded in a registry at ingest, replacing an earlier hard-coded global assumption.
- **Legality flags that gate distribution.** Each source carries [redistributable](#) (may WHG re-host and hand out copies at all?) and [downloadable](#) (redistributable *and* within a volume cap). Where a download is withheld, a reason is recorded — [licence-restricted](#) or [volume-exceeds-cap](#) . One source is currently gated out of downloads entirely on licence grounds.
- **Structured credit.** Contributor roles follow the NISO CRediT taxonomy, so bespoke acknowledgements and negotiated conditions are recorded as structured, exportable metadata rather than free-text prose.

4. How this works on the live site

4.1 Collecting terms from contributors

In the new contribution flow (the WHG *Workbench*, currently in beta and superseding the older upload path), a contributor chooses a licence from the same controlled vocabulary the platform uses internally. A public reference page at whgazetteer.org/licenses — linked from the site's top menu — explains every licence in plain language, with an interactive picker, and links each to its official deed in up to 17 languages. **The same catalogue code drives both this page and the in-Workbench picker**, so a contributor never sees two different descriptions of the same licence.

To be precise about the transition: under the *legacy* upload path a contributor accepted a single blanket licence via a checkbox, and no per-dataset licence was recorded. That historical material therefore still needs a backfill, and the wording of that blanket acceptance is the text we ask you to review at §7.2. The per-source model described here is fully populated for the **authority gazetteers** that make up the 50.5 m indexed records (§5); extending it over the legacy contributed datasets is separate, tracked work.




Fig. 1 — /Licenses/, the public licence reference. Two things to note. The first callout states WHG's own overlay licence and says explicitly that it is asserted *in addition to* — *never instead of* — each source's own terms. The second introduces **"By arrangement"**: a contributor may signal that although their chosen licence does not itself grant commercial use or adaptation, they will consider requests case-by-case. WHG passes such enquiries on; the contributor's contact details stay private. This is our current mechanism for routing a commercial enquiry to a rights holder without WHG purporting to grant rights it does not hold.




Fig. 2 — the vocabulary, grouped from open to restrictive. Each licence is reduced to four flags a non-specialist can act on: *Commercial use*, *Adaptations*, *Attribution*, *ShareAlike*. These are the same flags the platform uses to gate distribution, so what a contributor is shown and what the system enforces are guaranteed to be the same thing.

Fig. 3 — the licence field in the contribution flow. Licence selection sits in the first section of the form, alongside title and description — before any data is entered, not buried at publication. The amber chip records the optional “commercial use by arrangement” condition described in Fig. 1.

Fig. 4 — the guided picker. Contributors who do not know which licence they want answer four plain questions (“Allow commercial use?”, “Allow adaptations?”, “Require attribution?”, “Require ShareAlike?”) and the list filters to the licences that match. It is the same component and the same catalogue as Fig. 2.

4.2 Surfacing terms to reusers

On the reuse side, licence information is attached to the data at every level at which a user can encounter it: the gazetteer list, the individual gazetteer, and the individual place record.




Fig. 5 — every source gazetteer carries a colour-coded licence badge. Green for attribution-only, amber for ShareAlike, red for non-commercial, grey **C** for bespoke/custom terms. A user choosing which sources to search sees the licensing consequences before selecting, not after.




Fig. 6 — per-gazetteer detail (*Index Villariz*). Record count, temporal extent, **rights holder** ("Stephen Gadd & Alexis Litvine (digital edition); original 1680 work by John Adams (public domain)"), a link to the source, a formatted citation with DOI, and a download button that is only offered where the source's flags permit redistribution. Note that rights are expressed as a *chain* — a modern licensed digital edition over a public-domain original — which a single licence field could not represent.




Fig. 7 — licensing at the individual record. Clicking any place on the map gives a popup whose footer reads *SOURCE & LICENCE — Index Villariz — CC-BY-SA-4.0*. The badge is tiered: it shows the SPDX identifier at a glance, expands on hover into the full licence name plus the rights-holder statement, and links through to the licence deed. A user who copies a single coordinate out of WHG can see, in two seconds, whose data it is and on what terms.

4.3 Machine-readable terms in the API — the remaining gap

The design requires that any API response spanning multiple sources carry an **aggregated statement of terms at the root of the returned document**, listing every source present in the results with its licence, so a programmatic consumer can comply without having to look anything up. Checking this against the live service while preparing this briefing showed it was only half-built; it has since been completed and deployed:

Surface	Carries source terms?	Status
Interactive UI (Figs. 5–7)	Yes — full licence, flags, rights holder, deed link	Done
Single-record API responses	Yes — an attribution object with licence and flags	Done
Standalone resolver endpoint	Yes — callers may ask "what are the terms for these namespaces?"	Done
Multi-record result sets (search, index, spatial)	Yes — aggregated terms block at the response root	Done
Reconciliation responses	Yes — root terms block, plus a source tag on every candidate	Done
The aggregated block itself	Yes — now carries the licence and its flags, not only names and citations	Done

This mattered most for **reconciliation**, which is arguably WHG's larger delivery channel and the one that blends candidates from many differently-licensed sources into a single response. A consumer previously received results spanning several licences with no machine-readable way to tell which was which — and therefore could not comply with per-source terms even when willing to. Every response now states, at its root, each contributing source with its licence and the two flags a reuser must act on (*permits commercial*

use, *share-alike*), and each reconciliation candidate is tagged with the source it came from. Where a source's licence is not recorded, the response reports *null* rather than defaulting to something permissive.

Two data-quality findings from the same exercise — these, not the plumbing, are the real gap.

(i) OpenHistoricalMap is recorded in our registry as CC0-1.0 . OpenHistoricalMap publishes under ODbL, as its sibling OpenStreetMap does (and as our own registry correctly records for OSM). This understates a share-alike obligation on roughly 0.9 million records.

(ii) Five authority gazetteers — 113,578 records — have a *rights holder* recorded but no licence at all: CHGIS, the Ancient Parishes of England & Wales, Native Land, the Digital Gazetteer of the Song Dynasty, and UK Historic Counties. For a licensing conversation this is the more serious category: a non-commercial licence can be cleared or excluded deterministically, whereas unrecorded terms cannot be reasoned about at all.

Both originate in the authority metadata at ingest and are being corrected there, so the fix propagates to every surface. Together they are why the per-namespace licence audit (§8) is the item we rank first: the machinery is now in place everywhere, and it is only ever as accurate as that table.

5. What the corpus actually contains

A recurring worry is that respecting upstream terms would gut any commercial offering. The measured composition says otherwise. Record counts below are live figures from the platform registry; licences are as currently recorded, pending completion of the per-namespace audit.

Source (namespace)	Records	Licence	Commercial use
OpenStreetMap osm	20,289,895	ODbL-1.0	✓ but share-alike
GeoNames gn	13,378,039	CC-BY-4.0	✓ attribution
Wikidata wd	11,459,393	CC0-1.0	✓ public domain
Getty TGN tgn	2,991,044	ODC-By-1.0	✓ attribution
GB1900 gb	1,174,449	CC-BY-SA-4.0	✓ but share-alike
OpenHistoricalMap ohm	905,205	CC0-1.0 (<i>under review</i> — see §4.3)	✓
China Historical GIS chgis	82,117	not recorded	✗ excluded
Trismegistos tm	64,196	CC-BY-SA-4.0	✓ but share-alike
Pleiades pl	25,561	CC-BY-3.0	✓ attribution
Index Villaris iv	24,000	CC-BY-SA-4.0	✓ but share-alike
Ancient Parishes of England & Wales kain_par	23,177	not recorded	✗ excluded — download gated
Alcedo alc	18,223	CC-BY-NC-4.0	✗ excluded
Ottoman NFS Gazetteer ofs	16,296	CC-BY-4.0	✓ attribution
Cliopatria clio	15,690	CC-BY-4.0	✓ attribution

Source (namespace)	Records	Licence	Commercial use
HGIS de las Indias hgis	14,152	CC-BY-SA-4.0	✓ but share-alike
PeriodO po	9,017	CC0-1.0	✓ public domain
Ottoman Gazetteer og	6,260	CC-BY-NC-4.0	✗ excluded
Native Land nl	4,364	not recorded	✗ excluded
Digital Gazetteer of the Song Dynasty dgsd	3,828	not recorded	✗ excluded
D-PLACE dp	2,599	CC-BY-NC-4.0	✗ excluded
GBHGIS boundary sets (4 gazetteers)	2,850	CC-BY-SA-4.0	✓ but share-alike
UN Countries un	247	public domain	✓
UK Historic Counties ukhc	92	not recorded	✗ excluded

The headline figure: 99.7% of the indexed corpus carries a licence that permits commercial use. Of 50,510,694 records, **28.8 m (57.0%)** are attribution-only or public domain — free of any copyleft obligation; **21.6 m (42.7%)** are commercially usable but *share-alike*; only **27,082 records (0.05%)** are non-commercial; and **113,578 records (0.22%)** have **no licence recorded at all**.

So the commercially-consequential decision is *not* whether to exclude the non-commercial material — that is a rounding error by volume. It is **whether to include the share-alike sources**, which are 42.7% of the corpus and are dominated by a single source (OpenStreetMap under ODbL).

Three qualifications keep this honest. First, **the 0.22% with no recorded licence is the category that should concern us most**, and it is larger than the non-commercial one. A non-commercial licence is a known quantity: it can be honoured, excluded, or renegotiated. An unrecorded licence cannot be reasoned about at all — we can neither offer the data nor confidently refuse it. Correcting this is the first item in §8, and we would rather show it to you than round it into a "bespoke terms" bucket.

Second, **volume is not value**: the restricted and unrecorded records are specialist, hand-curated material (Chinese historical administrative units, English parish boundaries, Indigenous territories) whose scholarly value per record far exceeds a GeoNames row — and it is precisely the material a scholarly aggregator exists to provide. Third, **places are heavily duplicated across authorities**, so a commercial consumer denied a restricted source will usually still match the same place in GeoNames, Wikidata or TGN. The real loss is the tail of places that exist *only* in a restricted set.

What ShareAlike would require of a commercial licensee

Because share-alike is 42.7% of the corpus, its obligations are worth stating precisely. Both ODbL and CC BY-SA permit commercial use, but attach copyleft:

- CC BY-SA reaches *adaptations of the content itself* — any adaptation must be re-licensed under the same terms.
- ODbL attaches to the *database*, not necessarily the end product. A licensee who builds and publicly distributes a derivative database must offer that database under ODbL with alterations available — but an application that merely *displays* the data is not thereby forced open (ODbL's "Produced Work" carve-out).

The practical consequence for a commercial partner is that they generally cannot keep a derived place-database proprietary if they distribute it, and SA terms may be incompatible with other licences in their stack. By contrast CC BY, CC0 and ODC-By are attribution-only — the cleanest 57% for a commercial consumer, and the subset we would suggest as the default commercial offering.

6. What WHG could legitimately commercialise

We cannot sell what we do not own. On our reading, the WHG-owned or sublicensable layers are these — and we would value your view on whether the second is as defensible as we believe:

- **The software** — already released under BSD-3.
- **The reconciliation and linkage service.** This is the matching, ranking and cross-source crosswalk — the intellectual work WHG genuinely creates, distinct from the underlying records. It is what a partner such as a genealogy or mapping product actually wants, and it is the layer we could plausibly put behind a commercial licence *even where we cannot license the source records themselves*. This is the core of our commercial thinking.
- **Whole contributed gazetteers whose contributors granted commercial rights** — enumerable from the per-source flags, and now including those flagged "commercial use by arrangement" (Fig. 1), who have signalled willingness to be approached.
- **WHG's own editorial content and curation** — currently CC BY-NC-4.0.

A dual-licence structure (free for non-commercial use, negotiated licence for commercial) is the model we have in mind, built over the service layer rather than the aggregated records. We are aware of comparable arrangements in the research-data sector and would welcome guidance on how Pitt would prefer such a structure to be framed.

7. Specific matters on which we seek guidance

7.1 A bespoke co-authorship condition

One source (a set of English parish transcripts from the University of Portsmouth) comes with terms asserting that (a) Portsmouth holds some copyright and the files are "definitely not public domain"; (b) the depositor is content to release them as CC BY-NC-SA; and (c) that "BY" means he *should be included as a co-author* on any publication making substantial use of the data.

Our reading: (a) is broadly sound though not quite as stated — the raw facts are not copyrightable, but the *compilation* is, and under UK law the **sui generis database right** protects exactly the substantial investment described (OCR-scraping software plus extensive manual clean-up). Because that work was done by university staff and students, the rights holder is most likely the **institution rather than the individual**, so any commercial use would need institutional sign-off.

On (c) we believe the claim over-reaches. Under Creative Commons, the Attribution requirement means *credit* — naming the creator, linking source and licence, noting changes. It does not confer or require co-authorship. Attribution is not authorship: co-authorship follows from a substantive intellectual contribution to the specific new work, or from contract — not from a licence's attribution clause. A depositor may legitimately *negotiate* co-authorship as a separate bespoke condition of access, but as a bespoke term it binds only WHG and cannot be passed through to downstream users the way a licence can.

Why it matters operationally: a mandatory co-authorship term is in practice *more* restrictive than NC — it cannot be satisfied automatically, does not scale, and WHG cannot honour it downstream. This source is consequently the one currently gated out of downloads entirely. *We would welcome confirmation that our reading of attribution vs. authorship is correct, and guidance on how to negotiate such terms in future deposits.*

7.2 The contributor grant wording

Uploaders accept a licence grant at the point of contribution. This is the most legally load-bearing text on the platform: it determines what WHG may subsequently do with contributed data, and CC BY-4.0 in particular is **irrevocable**, explicitly permits commercial use, and forbids the licensor from layering additional restrictions on it. We would like this specific wording reviewed and signed off, and guidance on whether we should be asking for a licence, a broader sublicensable grant, or both.

7.3 Is "non-commercial" a mechanism worth relying on?

Our internal view — shared by colleagues — is that NC is a weak instrument: it is poorly defined, catches uses we would want to permit (commercial academic publishers, tuition-charging universities), and does not reliably prevent the uses we would want to stop. We currently apply CC BY-NC-4.0 as WHG's own overlay licence. *We would value a view on whether that remains the right default for the WHG layer, or whether a different instrument would serve better.*

7.4 A declared-use-case attestation

Per-source flags only become enforceable once we know what a user intends to do. Our proposal is to capture a **declared use case** at the point of access — at download, and once at API-token issuance rather than per request — **defaulting to non-commercial**, with commercial use as an explicit opt-in. The declaration would then combine with the per-source flags: the default path receives the full set under the appropriate terms; a commercial declaration is served only the commercially-licensed subset and routed to a licensing enquiry rather than a flat refusal.

Our belief is that this creates an **attestation record** evidencing good-faith enforcement of upstream terms and places compliance responsibility on the declaring user — a meaningful shield for a platform redistributing third-party data — while doubling as a commercial-enquiry funnel. *We would value your view on whether that protection is real, and on what the declaration would need to say to be worth anything.*

8. Remaining work, priority-ordered

#	Area	Change	Priority
1	Per-namespace licence audit	Two known defects to correct at source, plus verification of the rest: OpenHistoricalMap recorded as CC0 rather than ODbL (~0.9 m records), and five gazetteers (113,578 records) with a rights holder but no licence (§4.3). Every downstream surface is now wired, and is only as accurate as this table.	Highest
2	DOI registration metadata	DOIs minted for WHG datasets currently register a hard-coded CC-BY-NC-4.0 rights statement with DataCite regardless of the dataset's actual licence. This has the longest tail of any item — DOI metadata is externally cached and hard to walk back — and should be fixed before the next minting run.	Highest
3	Contributor grant wording	Legally load-bearing (§7.2). Decide deliberately what grant we ask for; counsel signs off on this text specifically.	Highest
4	Reconciliation & result-set API responses	Aggregated terms block at the response root of search, index, spatial and reconciliation, plus a source tag on every reconciliation candidate (§4.3). Enables consumer	Done

#	Area	Change	Priority
		compliance and lets WHG gate sources whose terms cannot be honoured.	
5	Aggregated attribution block	Extended to carry the licence and its flags, not only source name, citation and record count.	Done
6	Declared use case & gating	Implement the attestation described in §7.4, subject to guidance from this meeting.	Med
7	Legacy contributed datasets	Datasets uploaded through the older path carry no per-dataset licence (they were covered by a blanket acceptance checkbox). Backfill once the §7.2 wording is settled, so the per-source model covers contributed data as well as the authority gazetteers.	Med
8	Legacy static pages & download artefacts	A small number of older templates, download README files and email texts still carry the pre-audit blanket licence statement. Text-only changes; to be aligned once the wording in §7.2–7.3 is settled.	Med
9	Licence vocabulary, per-source capture, contributor picker, UI badges, API terms block	Built and in production (§3, §4). The machinery is complete; what remains is data accuracy and coverage.	Done

The sequencing principle is that the *wording* decisions (items 2, 3, 6) are legal ones and gate the technical work; once settled, the single-source-of-truth architecture already in place makes every surface consistent at once.

This document is a technical and editorial briefing prepared to inform discussion. It identifies factual positions and structural logic; it is not legal advice. Statements about copyright, database right and Creative Commons interpretation are offered in good faith to frame the questions, not to resolve them. Decisions on contributor grant wording, commercial licensing, third-party rights, and any change to DOI metadata should be confirmed with the Office of Innovation and legal counsel. Record counts and recorded licences are live figures taken from the platform on 29 July 2026, pending completion of the per-namespace licence audit described in §8.